



Fake News Detector

Using ML to Help Tackle Misinformation

Jesús Badenes @ Ironhack



What will we cover today?

01

Why Solve for This?

02

The Approach

03

The Data

04

The Models

05

The Application

06

Final Reflections





01

Why Solve for This?



The State of Affairs

Creating and Distributing Fake News Has Never Been Easier


- **AI** tools produce text **massively**.
- **Social media** amplifies **misinformation** and is a growing source.

Trust in Media Has Never Been Lower

- Ipsos* reports a **5% global drop in media trust** over 5Y.
- Reasons: **fake news** spread and doubts about **media integrity**.

Good Journalism is at the Heart of Democracy

- Research links **healthy media** to **effective democracies**.
- **Media trust** can boost **voter turnout** by **13%**.



Why is Identifying Fake News so Difficult?



Information $\times 10^n$ = No Information


- The **Internet** offers **vast information**, but humans struggle to process it all.
- Selecting reliable information is **effort-intensive** and **time-consuming**.

Algorithmic Echo Chambers

- We manage information through **quality news** (often costly) and **algorithmic curation**.
- Social network algorithms often prioritize **attention over truth**, showing us preferred (even if untrue) content.

Confirmation Bias

- People naturally favor information that **aligns with their beliefs**.
- This "**Confirmation Bias**" makes detecting fake news hard.





02

The Approach






Identifying Fake News at Scale Through ML

Using a trained model, we can quickly help identify potential fake news

Dataset of Fake and True News



Data
Cleanup

Binary Classification Algorithms

- Logistic Regression
- Decision Tree Classifier
- Random Forest Classifier
- Support Vector Machine
- Feedforward Neural Network
- Recurrent Neural Network

Model
Deployment

User Application

**Deploy Model for
User Application**


Streamlit





03

The Data

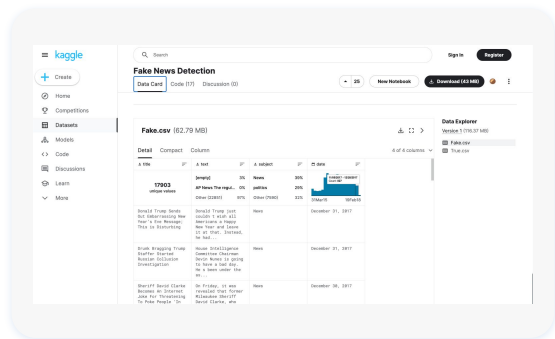




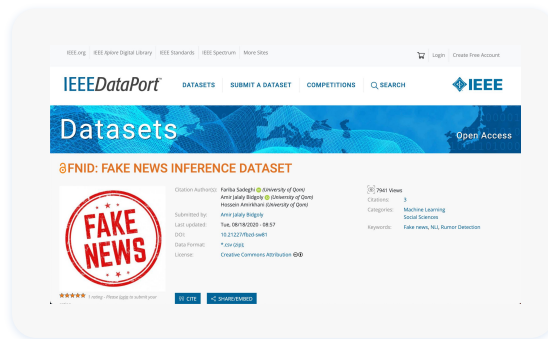
Data Sources

To train a binary classification model, we need a labeled & diverse dataset of news articles

Kaggle Dataset (40k)



IEEE Dataset (20k)



Modeled Features

- ❖ News Title
- ❖ News Text

Challenges:

- It is essential to train our models on previously classed data from fact checking organizations.
- We expanded our dataset to two sources due to overfitting concerns.



Data Cleaning



Removing noise boosts model accuracy, ensures consistency, minimizes sparsity for better efficiency, and helps in reducing biases, ultimately promoting better adaptability to new data

What Key Actions Did We Take?


- **Lowercasing:** Convert all text characters to **lowercase** for consistency.
- **Removing Brackets:** Eliminate **square brackets** and their enclosed contents.
- **Non-Word Characters:** Replace **non-word** characters with spaces.
- **URL Removal:** Strip out **URLs** and web addresses.
- **HTML Tag Removal:** Erase **HTML tags** from the text.
- **Punctuation Removal:** Discard **punctuation** marks for cleaner text.
- **Newline Character Removal:** Erase **newline characters** for continuous text.
- **Alphanumeric Word Removal:** Strip words containing **digits**.
- **Expanding Contractions:** Convert **contractions** (e.g., "it's" to "it is").
- **Stopword Removal & Lemmatization:** Remove **common words** (e.g., "is", "an") and reduce words to their root form using **lemmatization**.



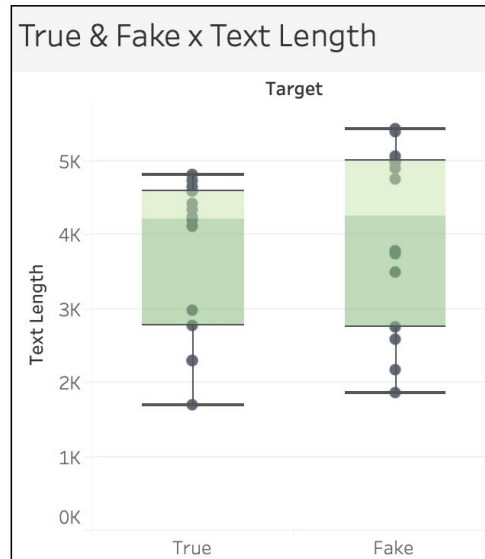


Exploratory Data Analysis (1/3)

Understanding the distribution of our data and analyzing patterns



Worked with balanced data



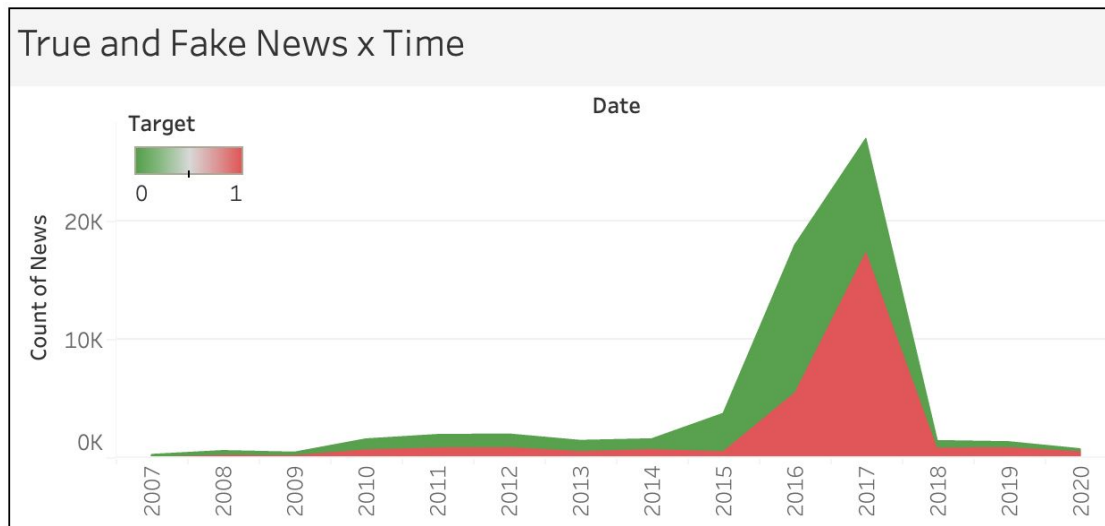
Observed that fake news were slightly longer





Exploratory Data Analysis (2/3)

Visualising a time series of news articles in the training set



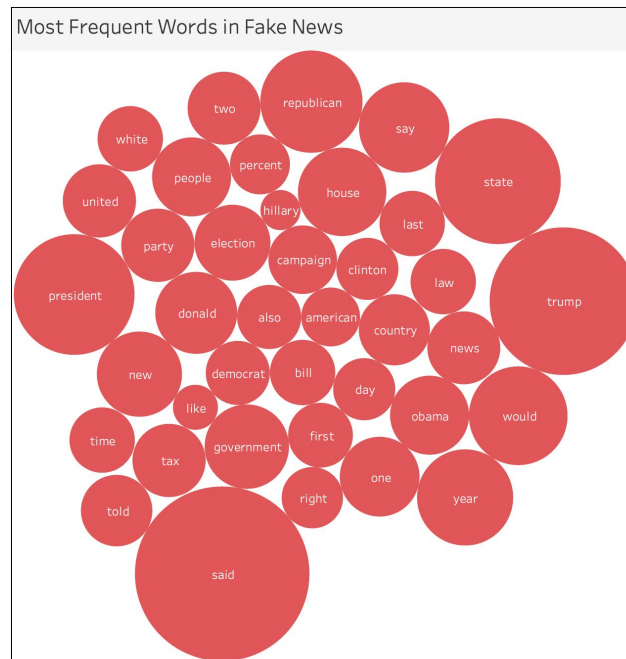
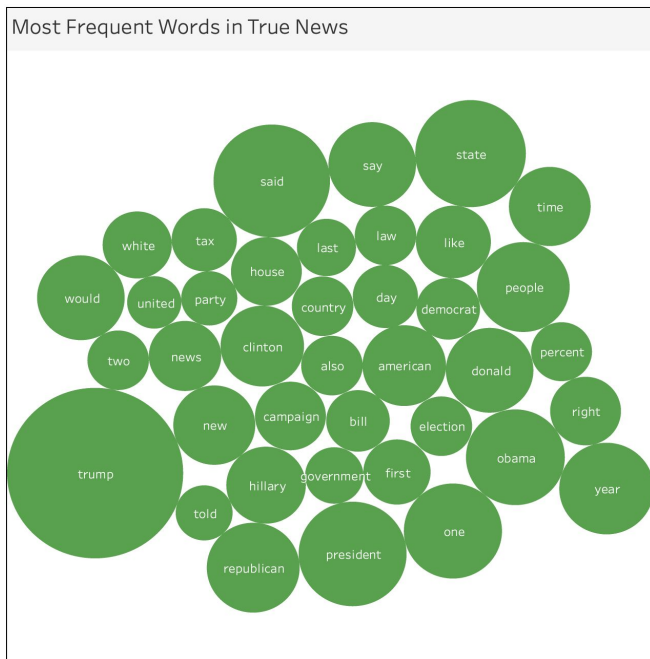
A majority of the articles in the training set are from 2015 - 2018, and they are all in the range of 2007 - 2020





Exploratory Data Analysis (3 / 3)

Visualising most common words across true and fake news



We can observe a significant overlap in most common words, but there are key differences





04

The Models



The Models

After cleaning the data, we ran 6 models well fit for classification tasks

Model Results

Model	Accuracy
Logistic Regression	0.89
Decision Tree Classifier	0.85
Random Forest Classifier	0.88
Support Vector Machine	0.89
Feedforward Neural Net	0.84
Recurrent Neural Net	0.90

Key Observations

RNN Excellence: With a 0.90 accuracy, the Recurrent Neural Net outperforms all, showing its strength in capturing sequential patterns.

Classic Algorithms Shine: Logistic Regression and SVM both achieve an impressive 0.89 accuracy, highlighting their competitive performance.

Robust Data: Most models hover around upper 80% accuracy, indicating consistent and clean data.

Ensemble Benefit: Random Forest (0.88) outperforms its single counterpart, the Decision Tree (0.85), underscoring ensemble methods' advantage.


Simplicity vs. Structure: The RNN's superior performance over the Feedforward Neural Net (0.90 vs. 0.84) suggests the importance of data structure.





The Selected Model: Recurrent Neural Network

Our final model is a recurrent neural network with an 11-layer architecture





05

The Application





Empowering Users to Identify Fake News

Creating a Streamlit App for users to paste news articles and assess their credibility

Predicting True News...

Fake News Detector

News Title

President Donald Trump questioned the politics

News Text

President Donald Trump questioned the politics of his decision to finally acknowledge former President Barack Obama was born in the US, which he did late during the campaign in 2016, according to a source close to the White House.

The source said that shortly after he made the statement, Trump told aides that he would have done better in the polls had he continued to stand his ground on the birth certificate issue.

Trump has continued to question the legitimacy of Obama's birth certificate during private conversations in recent months, The New York Times reported Tuesday, citing

Predict

This news is likely true.

The model predicts a probability of 87.87 % that this news is fake.

Disclaimer: This prediction is based on a Deep Learning model called LSTM (Long Short-Term Memory). It is based on a variety of recurrent neural networks (RNNs) that are capable of learning long-term dependencies, especially in sequence prediction problems. This is a research project and as such, it may be prone to errors. We recommend doing further research before determining whether a piece of news is true or fake.

... And Fake News

Fake News Detector

News Title

The Pope's Coat Is Here to Ruin Your Faith

News Text

Future generations will acknowledge the vibe shift. It happened last weekend, when all of a sudden social media feeds filled with images of Pope Francis, typically a pious and plain dude, looking like a boss in a sleek white puffer coat. It was instantly a meme, a LOL in a sea of bad news. It was also not real. Someone created the image using the artificial intelligence tool Midjourney. But it fooled a lot of people—so many that news outlets began calling it “one of the first instances of wide-scale misinformation stemming from artificial intelligence.”

Just typing that sentence feels haunting. Like the first time you see someone in a red cloak in The Handmaid's Tale. Not that this

Predict

This news is likely fake.

The model predicts a probability of 87.34 % that this news is fake.

Disclaimer: This prediction is based on a Deep Learning model called LSTM (Long Short-Term Memory). It is based on a variety of recurrent neural networks (RNNs) that are capable of learning long-term dependencies, especially in sequence prediction problems. This is a research project and as such, it may be prone to errors. We recommend doing further research before determining whether a piece of news is true or fake.





06

Final Reflections





Challenges and Future Developments

Challenges

- **Data Diversity:** If the training data isn't diverse enough, the model might become too specialized, hindering its performance on unseen data.
- **Overfitting Risk:** While the model includes dropout and kernel regularization to combat overfitting, there's always a risk, especially with deeper networks.
- **Model Interpretability:** LSTMs, like other deep learning models, are often considered "black boxes", making it challenging to understand and explain their decision-making processes.
- **Scalability:** As the dataset grows, the computational cost of training such a complex model will increase.
- **Research & Confirmation Bias:** Relying exclusively on predictions to identify fake news can be risky and the need for the user to do further research is a limitation.

Future Developments

- **Expanding Data:** Adding more data sources will increase the model's ability to generalize.
- **BERT and Variants:** Leveraging models like BERT, RoBERTa, or DistilBERT which have been pre-trained on massive corpora can bring in the advantage of extensive prior knowledge.





Thank you